


Kubernetes Security with Calico and Open Policy Agent

Kevin Harris @kevbhar & Ray Kao @raykao

Open Source Solution Leads

Azure Global Black Belt

Microsoft Canada



Network policy

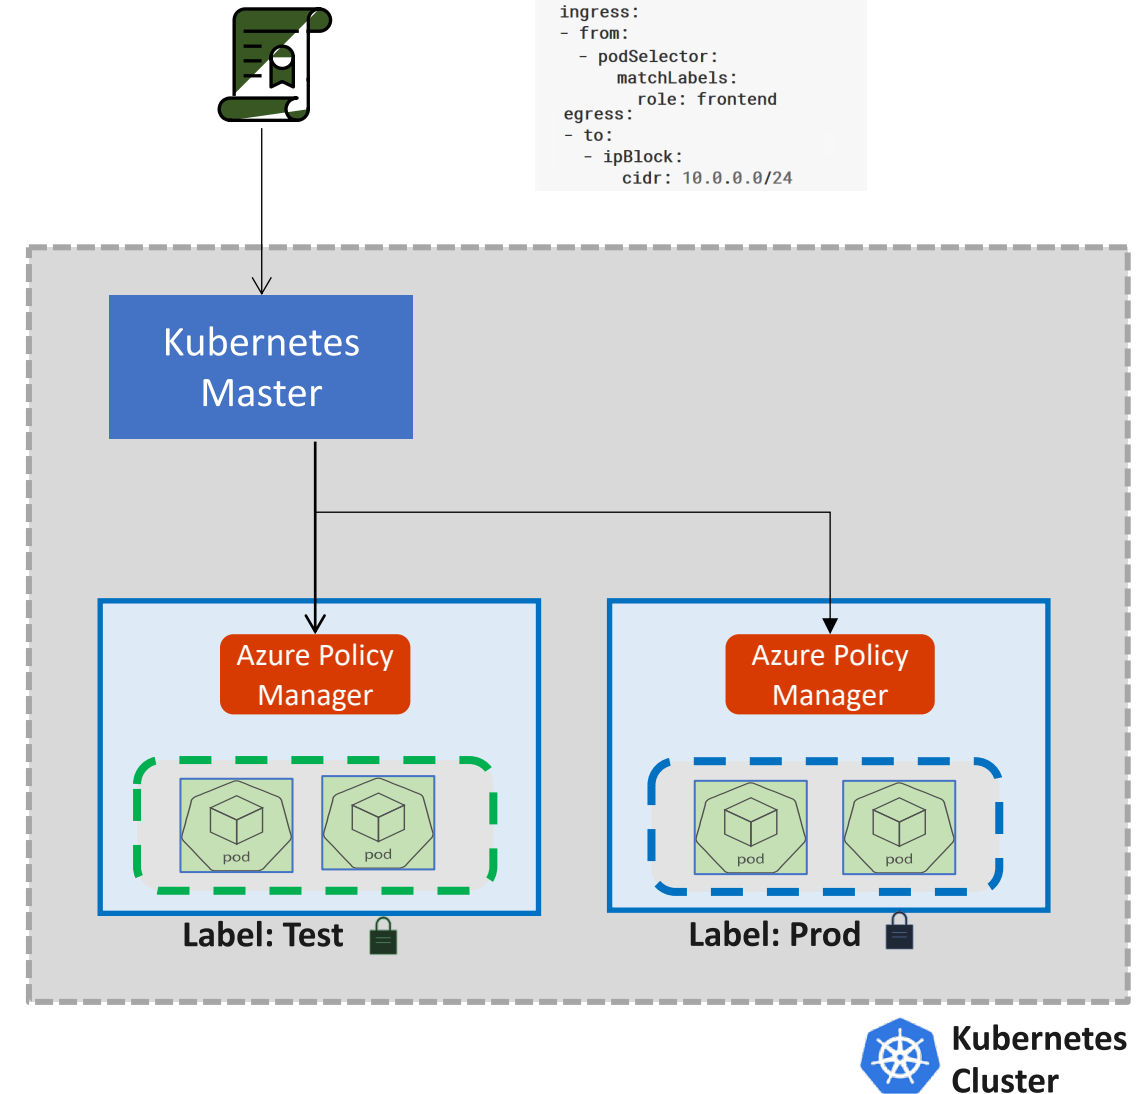
- All Pods are non-isolated by default
- Flat network. All pods can talk to other pods
- Accept traffic from anyone
- Multi stage/zone project this could expose security risks
 - 3 tier webapp.
 - Front end could technically talk directly to DB tier

Azure Kubernetes Network Policies

- Provides micro-segmentation for containers – like NSGs for VMs
- Label-based selection of Pods
- Policy resource yaml file specifies Ingress and Egress policies
 - Policies defined for a label
- Works in conjunction with Azure CNI
- Supports Linux hosts
- Supported in *aks-engine*
 - Set `networkPolicy` setting to `azure` in cluster definition file
- Supported on AKS in **Preview**

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
  - Ingress
  - Egress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          role: frontend
  egress:
  - to:
    - ipBlock:
        cidr: 10.0.0.0/24
```

`kubectl apply -f policy.yaml`



Network policy

- Pod Selector
- PolicyTypes
 - Ingress, egress

Ingress

- namespaceSelector
- podSelector
- ipBlock

```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: test-network-policy
  namespace: default
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
    - Ingress
    - Egress
  ingress:
    - from:
        - ipBlock:
            cidr: 172.17.0.0/16
            except:
              - 172.17.1.0/24
        - namespaceSelector:
            matchLabels:
              project: myproject
        - podSelector:
            matchLabels:
              role: frontend
      ports:
        - protocol: TCP
          port: 6379
  egress:
    - to:
        - ipBlock:
            cidr: 10.0.0.0/24
      ports:
        - protocol: TCP
          port: 5978
```

Network policy

- Recommend to create default policies that apply to all pods
 - Default deny all ingress traffic
 - Default deny all egress traffic.
 - Allow Specific Traffic (Ingress & Egress)

Calico Global Network Policy

```
apiVersion: projectcalico.org/v3
kind: GlobalNetworkPolicy
metadata:
  name: allow-tcp-6379
spec:
  selector: role == 'database'
  types:
    - Ingress
    - Egress
  ingress:
    - action: Allow
      protocol: TCP
      source:
        selector: role == 'frontend'
      destination:
        ports:
          - 6379
  egress:
    - action: Allow
```

Calico Global Network Set

```
apiVersion:
projectcalico.org/v3
kind: GlobalNetworkSet
metadata:
  name: a-name-for-the-set
  labels:
    role: external-database
spec:
  nets:
    - 198.51.100.0/28
    - 203.0.113.0/24
```

```
...
ingress:
- from:
  - namespaceSelector:
      matchLabels:
        user: alice
    podSelector:
      matchLabels:
        role: client
...
```

<- Single Rule where two conditions must be met

```
...
ingress:
- from:
  - namespaceSelector:
      matchLabels:
        user: alice
  - podSelector:
      matchLabels:
        role: client
...
```

<- two conditions either/or must be met

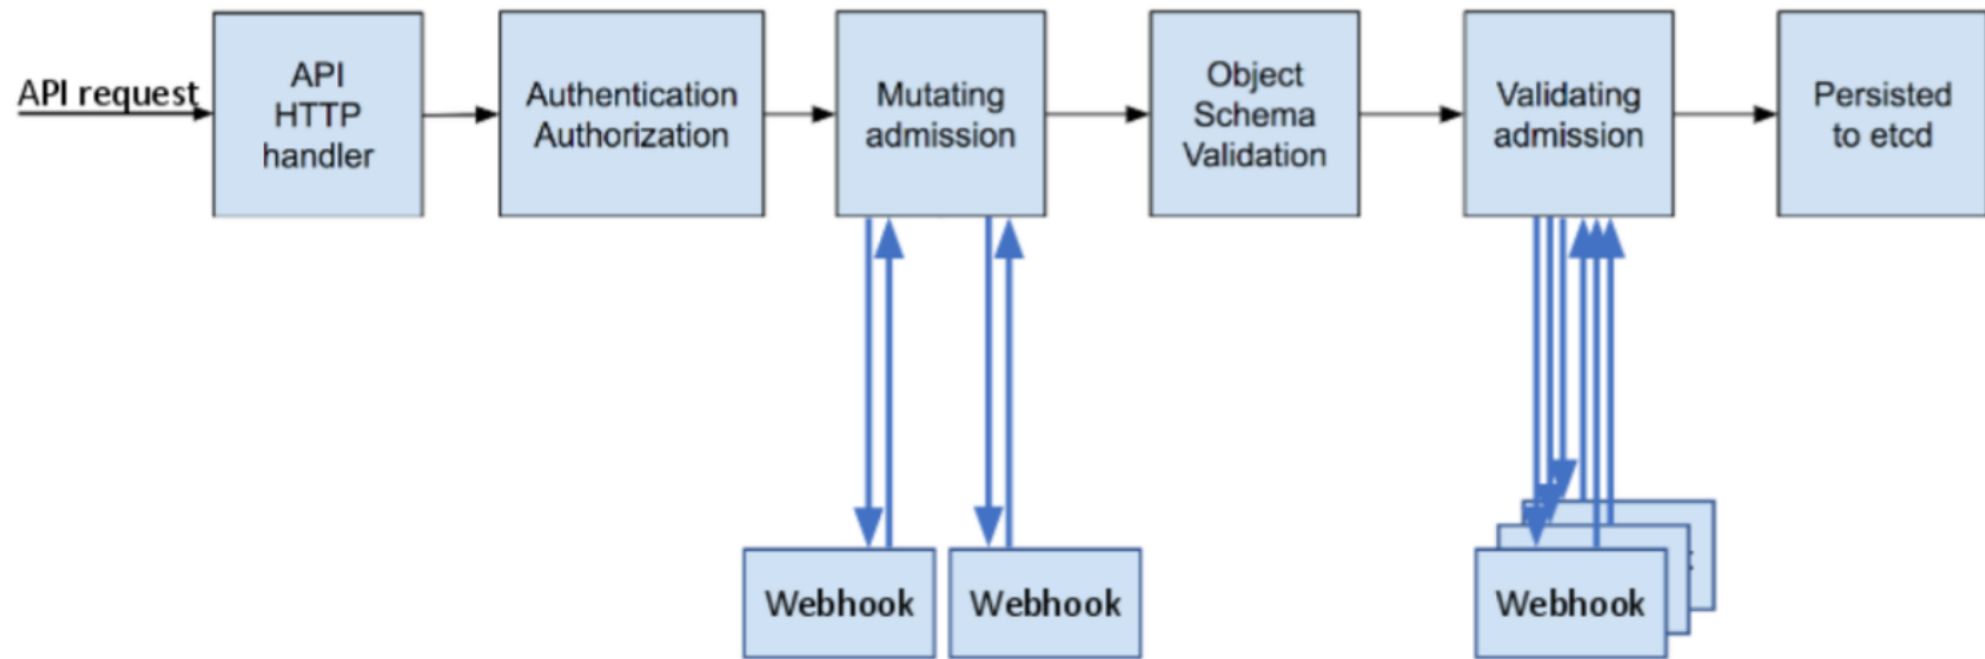
When to use Azure NSG vs. Kubernetes Network policy

- Azure Network Security Groups
 - Use it to filter North-South traffic, that is, traffic entering and leaving your cluster subnet
- Kubernetes Network Policies
 - Use it to filter East-West traffic, that is, traffic between pods in your cluster

Network Policy Demo

- Fruit Categorization App
- Web Frontend
- ML Backend for image recognition
- Signal R as a managed websocket service between FE/BE
- Default allow
- Create a policy (deny all by default)
- Open a policy to communicate between FE/BE services

Open Policy Agent



Admission Controller Phases

Dynamic Admission Control

- Validating Webhook
 - Allows you to intercept and validate requests
 - Can be run in parallel, as they don't mutate objects
 - Example use case: restricting resource creation
- Mutating Webhook
 - Executes the mutation by sending requests to webhook server
 - Matching webhooks are called in serial
 - Example use case: injecting side cars
- Policy Enforcement
 - Admission Control is policy based on Kubernetes objects.
 - Network Policy and PodSecurity Policy focus on data plane policy
 - RBAC is policy enforced on the user

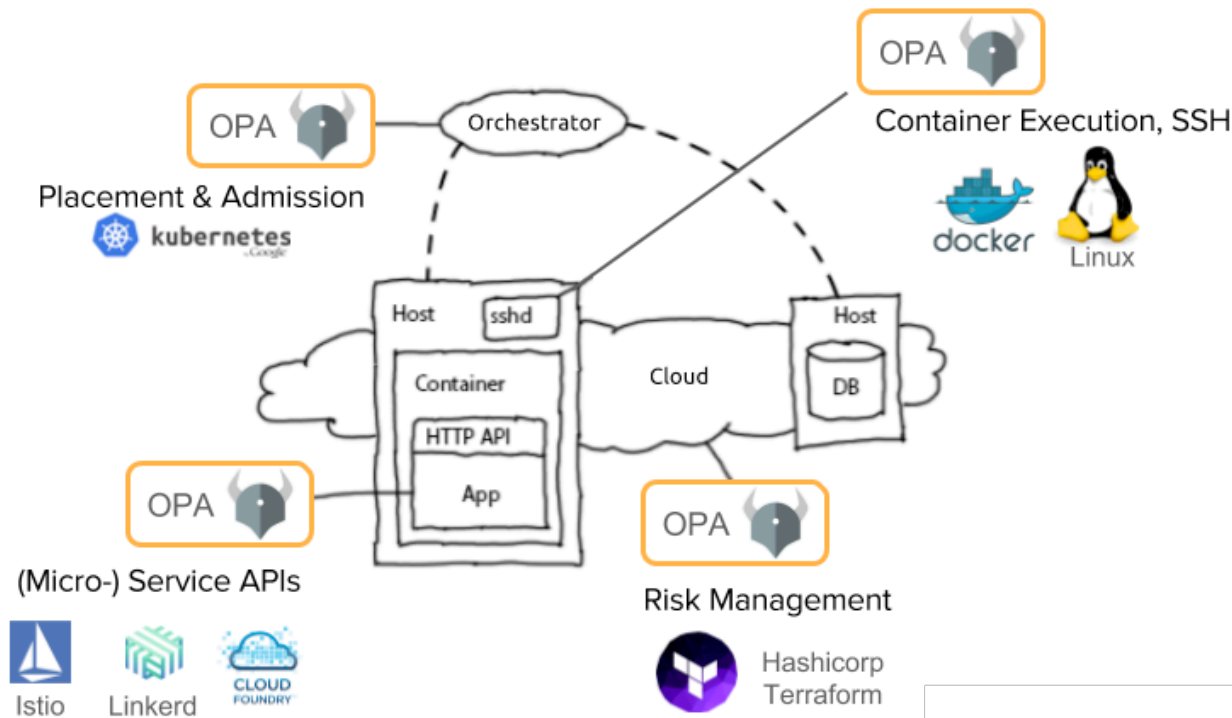
That's awesome! But...



Sample Admission Webhook

```
181
182     http.HandleFunc("/services", serveServices)
183     http.HandleFunc("/mutating-services", serveMutateServices)
184     http.HandleFunc("/healthz", serveHealthz)
185     clientset := getClient()
186     server := &http.Server{
187         Addr:      fmt.Sprintf(":%s", Options.PortNumber),
188         TLSConfig: configTLS(clientset, &certKey),
189     }
190
191     glog.V(2).Infof("starting webserver on port %s", Options.PortNumber)
192     glog.V(2).Infof("service annotation to match/mutate: %s: %s", Options.ServiceAnnotationKey, Options.ServiceAnnotationV
193
194     if err := server.ListenAndServeTLS("", ""); err != nil {
195         glog.Fatal(err)
196     }
197
198 }
```

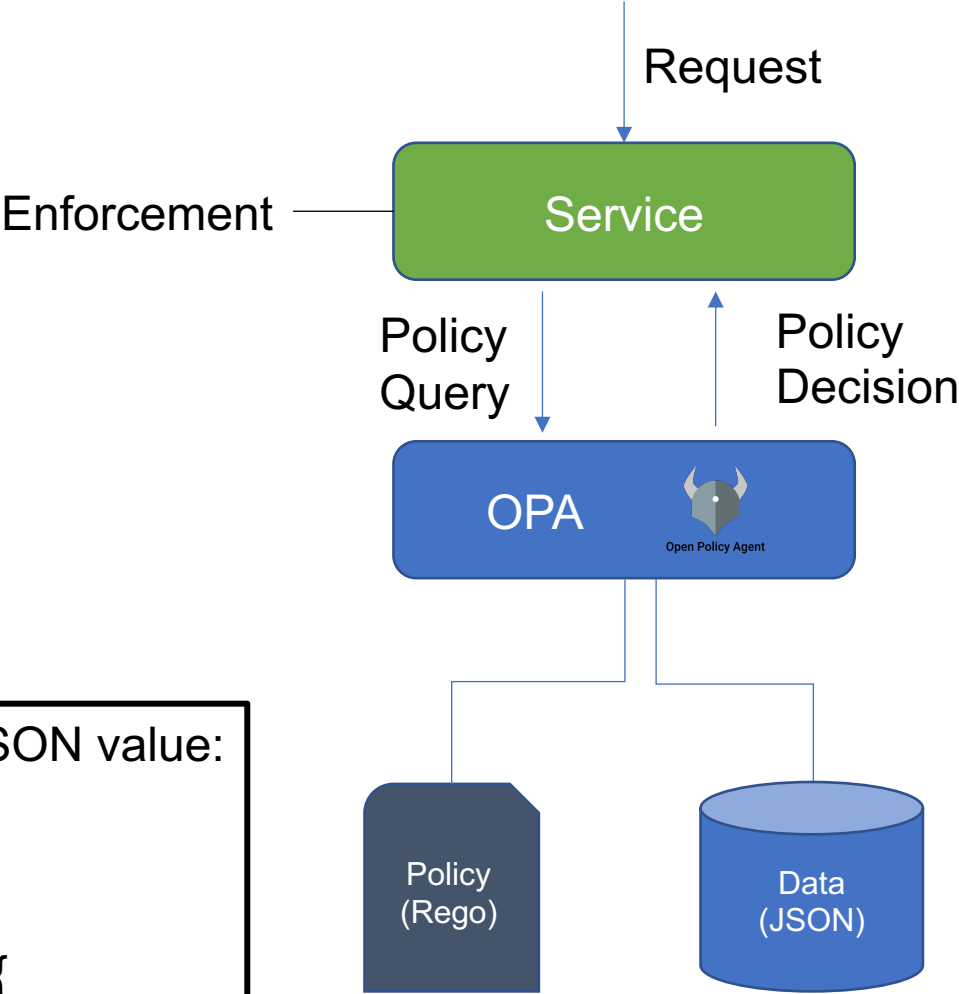
Open Policy Agent



- CNCF Hosted Sandbox Project
- General purpose policy engine
- Can be used across the stack
- Declarative policy language (Rego)



Open Policy Agent



Service refers to:

- Kubernetes API
- Custom API
- SSH Daemon
- Terraform
- Authorization API

Input can be any JSON value:

```
"kind": "Service",
  "metadata": {
    "annotations": {
      department: dev
    }
  }
```

Output can be any JSON value:

```
"true"
  "request annotated"
  " " "annotations": {
    costCenter: 8000
  }
```

Example Rego Policy

- Rego is a policy language and not a programming language, so don't think about sockets, methods, binary trees, etc.
- Think about two things: Logic and Data
- Rego logic is all queries. A query finds values for variables that make boolean conditions true.
- You write logic to search and combine JSON/YAML data from different sources.

```
deny[{
```

```
  "id": "conditional-annotation",
```

```
  "resource": {"kind": kind, "namespace": namespace, "name": name},
```

```
  "resolution": {"patches": p, "message": "conditional annotation"}, }] {
```

```
  matches[[kind, namespace, name, matched_object]] matched_object.metadata.annotations["Mr-T"]
```

```
  p = [{"op": "add", "path": "/metadata/annotations/cost-center", "value": "A-Team"}] }
```

Who manages all this policy?

Developer



Ice Kube

Deploys Apps

Platform Operator



Acid Burn

Creates And Maintains

OPA Policy

Audits Platform



The Governor

Kubernetes Policy Controller

- Kubernetes Policy Controller
 - Moving to OPA org, as a standard Kubernetes Policy Controller
 - Authorization module makes it possible to implement a blacklist in front of RBAC
 - Provides auditing features
 - Deployment consist of three containers: OPA, kube-mgmt., and Controller
- Examples:
 - Whitelist / blacklist registries.
 - Not allow conflicting hosts for ingresses.
 - Label objects based on a user from a department.
 - Block kubectl exec <pod>

The Good, The Bad, and Gotchas

- *Good*
 - OPA approach allows you to decouple policy from your applications
 - General purpose, so can be used outside of Kubernetes context.
- *Bad*
 - There can be a learning curve to Rego.
 - Can cause latency, but's negligible for most apps. (more of a consideration)
- *Gotchas*
 - Mutating objects need to be handled with care. They can cause unexpected behavior to what the end-user expects.

Takeaways

- Focus on security is a ***must*** in any Kubernetes deployment.
- *Help **educate*** Security Teams on how to extend Kubernetes to integrate custom policies.
- Treat the Kubernetes cluster as ***immutable***, just like you do with applications.
- Multiple ways to accomplish policy
 - **Build all your own logic** and utilize dynamic admission control
 - **Utilize Open Policy Agent to simplify** deployment and logic for rule sets.

Example Policies

- <https://github.com/open-policy-agent/contrib>